



University
of Dundee



Shenyang University
of Technology



UNIVERSITY OF
BIRMINGHAM

Paper ID 124973

A Prompt Is Not Ground Truth:

Human Validation of Local Target States in Generated Images

Haorui Yu

University of Dundee
Dundee, United Kingdom

Hanwen Chen

Shenyang University of Technology
Shenyang, China

Qiufeng Yi*

University of Birmingham
Birmingham, United Kingdom

Why the image must be checked

A prompt specifies what should be generated. If the local target is omitted or ambiguous, the request alone does not establish what is visible in the image.

Observe first, then compare

200 images · 8 target families · 3 raters
120 target-present + 80 target-absent requests
25 images per family (15 present + 10 absent)

1 Inspect the generated image

Image + named target + anomaly family.
Full prompt and requested state are hidden.

2 Lock the individual responses

Each rater selects present, absent or uncertain.

3 Aggregate and compare

Majority vote is compared with the requested state.
Uncertainty and split votes remain unresolved.

Image evidence and disagreement

(a) Human-chair fusion

T134 | A4



Requested: present
R1 / R2 / R3: P / P / P
Majority: present
Match

(b) Text legibility

T29 | A2



Requested: present
R1 / R2 / R3: U / P / A
Majority: none
No majority

P: present A: absent U: uncertain

Confirmation of requested states

87.0%

174 / 200 local labels
confirmed

Wilson 95% CI: 81.6–91.0%

15

contradicted
(7.5%)

11

unresolved
(5.5%)


Unresolved: 5 uncertain majorities + 6 three-way splits.

92.1% among definite majorities (174 / 189)

Wilson 95% CI: 87.3–95.1%

Fleiss κ = 0.675 · Gwet AC1 = 0.757

Confirmation varied by target family



Descriptive comparisons; 25 items each; Wilson 95% CIs.
Count/repetition: 15/25 confirmed, 6 mismatches, 4 unresolved.

Implication for human control

Generation fidelity: image vs. human request.
Interpretation fidelity: AI claim vs. image evidence.

This study provides a local human reference. AI-judge accuracy and improved control after feedback remain untested.

Limits: 25 images per family; all raters were authors;
no stable generator ID or complete selection log.
Rates apply to this diagnostic collection.